

# DSP Implementation Notes

ESE2014

# fixed-point versus floating-point math

- when using software tools like Matlab or Octave, we are using floating-point number representations and seldom do we concern ourselves with issues like overflow or precision, since floating-point is a powerful representation:
  - it has a large range (due to the exponent portion), representing numbers spanning several orders of magnitude
  - it has excellent precision, especially on desktop computers, where many bytes can be used to represent the mantissa
- in embedded systems, while floating-point units are available on some processors, fixed-point arithmetic is still important:
  - a fixed-point implementation can be more portable (since it works on simpler MCUs without floating point units)
  - it can be faster or more energy efficient to implement arithmetic in fixed point, even if a floating-point unit is available
  - because of limited range and precision, we must be careful in our use of fixed-point techniques

# fixed-point math

- in fixed-point arithmetic, we are actually using integers in calculations, keeping in mind an explicit scaling factor:

$$x = x^{\text{norm}} \times 2^M$$

↑  
a normalized quantity, like the variables in Matlab, on some range like [-1,1]

↑  
the scaling factor, where  $M$  represents the number of bits used in the fractional part of the fixed point number,  $x$

- note that if  $x^{\text{norm}} = 1.0$ , then  $x$ , in binary, looks like:

0000 0000 0000 0001 . 0000 0000 0000 0000 ← 32-bit word

31 30 29 28 27 26 25 24 23 22 21 20 19 18 17 16 15 14 13 12 11 10 9 8 7 6 5 4 3 2 1 0

↑  
implicit “binary point”

# fixed-point examples

- $M = 16$ ,  $x^{\text{norm}} = 0.536$ ,  $y^{\text{norm}} = 0.251$
- in binary form,  $x^{\text{norm}}$  is: 0000 0000 0000 0000. 1000 1001 0011 0111
- we get:
  - $x = 35127$
  - $y = 16450$
- addition/subtraction is straightforward:
  - $x + y = 51577 \Rightarrow$  a “normalized” value of  $51577/2^{16} = 0.787$  (correct)
- multiplication is more tricky, because you must keep in mind the scaling factors:
  - $x * y = 577\,839\,150 \Rightarrow$  a 30-bit quantity (running the risk of overflow!)
    - this is a normalized value of  $577\,839\,150/2^{32}$

# floating-point examples (cont'd)

- division can be awful: while we don't need to worry about the scaling factors (they cancel out), we can get substantial rounding errors:
  - $x/y = 2.1354$ , which gets “integerized” to 2 (incorrect!) !
  - to deal with divisional rounding, we can use a larger range for our dividend,  $x$ , for example,  $[-10, 10]$  instead of  $[-1, 1]$ , however, with larger numbers, we push the risk of overflow on multiplications
  - you have to decide what's more important for you: avoiding overflow in multiplications, or rounding errors from divisions
- fortunately, in most filtering applications, we are dealing primarily with addition (of signals) and multiplications (signals with filter coefficients or amplifier gains).
- if it all possible, it is best to have power-of-2 divisors ( $y$ ) with integers, as this results in a simple bitwise-shift operation, which can be implemented with great speed on any MCU

## floating-point examples (cont'd)

- in the previous division example,  $y = 0.251$  which approximately  $0.250$ , or  $2^{-2}$ ; therefore  $x/y$  is approximately  $x \cdot 2^2$ , or simply shift  $x$  twice to the left.

This results in:

- $x \ll 2 = 0000\ 0000\ 0000\ 0010 . 0010\ 0100\ 1101\ 1100$

$$= 140508$$

$$= 2.1440$$

which is closer to  $2.1354$  than  $2$  is (error of  $0.0086$  versus  $0.1354$ ).

- a cool trick if you can make it happen!

# scaling filter coefficients

- whether fixed or floating point, we should strive to have the signals in our filter realizations vary roughly in the same range; for most purposes, we can think of our signals as ideally “swinging” in the  $[-1, 1]$  range
- by doing so:
  - we can more easily implement analog versions of our filters, since we want signals to vary within (and fully utilize) the supply limits in order to maintain the highest signal-to-noise ratios
  - we can reduce the possibility of underflow or overflow errors
- given a state-space realization of a filter, this can often mean “balancing the coefficients, so that they aren’t too many orders of magnitude apart