

FreeRTOS Queues & Mutexes


RTS-W5-1
FreeRTOS-6

- Please have students read corresponding chapter from Richard Barry's book

TASK 1 (writing)

QUEUE (BY COPY)

TASK 2 (reading)



by default, items are sent to back/tail of queue (FIFO)

* A QUEUE IS A FIFO DATA STRUCTURE. DATA IS READ FROM HEAD AND IS PLACED AT TAIL

API Functions

- QueueHandle_t x QueueCreate (port BaseType_t ux QueueLength, port BaseType_t ux ItemSize);
 ↑
 type names can change, depending on port
- BaseType_t x QueueSendToFront (QueueHandle_t x Queue, const void* pvItemToQueue, TickType_t xTicksToWait)
- BaseType_t x QueueSendToBack (QueueHandle_t x Queue, const void* pvItemToQueue, TickType_t xTicksToWait)
- BaseType_t x QueueReceive (QueueHandle_t x Queue, const void* pvBuffer, TickType_t xTicksToWait)

- UBaseType_t ux QueueMessagesWaiting (QueueHandle_t xQueue)

Mutex (Binary Semaphore)

- used for protecting a resource from access by more than one task at a time
- see: PRIORITY INVERSION
(when a task of lower priority is executing in spite of the fact that higher priority tasks are in the READY state)

→ EXERCISE:

- have students read Mutex-related material and compile their own lists of relevant API functions.

needed elements to use a mutex:

*include "semphr.h"

type: xSemaphoreHandle

API functions (students should find themselves)

- xSemaphoreTake()
- xSemaphoreGive()
- xSemaphoreCreateMutex()

Real-time Resources

RtS-W5-3
RTR-1

- A number of system resources must be sized and managed in any real-time embedded system, including:

PROCESSING: the number and type of cores available

MEMORY: all volatile and non-volatile storage elements

I/O: buses, digital communications (serial/parallel) used for sensor information decoding and encoding of actuator data
- includes interconnections between cores.

- traditionally, the focus of real-time resource analysis and theory has centred around processing and executing multiple services on single core.
- resources are allocated to specific threads of execution; the mechanics of preempting a running thread, saving its state, and dispatching a new thread is called a thread context switch.

Scheduling \longleftrightarrow implement a policy } how the RTOS makes a decision

preemption } \longleftrightarrow context-switching mechanisms } how the RTOS implements its policy
dispatch

- Some important factors affecting real-time resource allocation:

- SPEED or clock rate (GHz)
- EFFICIENCY or clocks per instruction (CPI)
or instructions per clock (IPC) \rightarrow relevant in pipelined execution

- Algorithm complexity:

C_i = instruction count on longest execution path for service i , and, ideally, deterministic (known a priori); if not known exactly, a WCET should be used (worst-case execution time).

- Frequency of Service Requests, F_i

$$T_i = \frac{1}{F_i} = \text{"service release period"}$$

- Latency Issues

- often associated with I/O, interconnected:

- arbitration (contention) latency for shared I/O devices
- read latency
- transit time from device to CPU core
- registers, tightly coupled memory (TCM), L1 cache for zero wait state
- single cycle access
- read/write latencies

L2 cache is memory on the main CPU chip

- Bandwidth

- average bytes or words transferred per unit time
- can be "uniform" or "bursty"

- Queue Depth

- queuing can keep the system running well, as long as they are not filled \rightarrow which leads to stalls...

- CPU Coupling

- DMA channels can decouple the CPU from I/O
- memory-mapped I/O strongly couples CPU to I/O (devices need more attention)

Memory Hierarchy (from least to most latency)

- Level-1 Cache
- Single cycle access
- Harvard Architecture (separates data and instructions)
- Level-2 Cache or TCM
- MMR (memory mapped registers)
- Main Memory - SRAM, SDRAM, DDR
- MMIO
- nonvolatile memory: Flash, EEPROM, NOVRAM

→ the way your memory is organized and how your devices are "memory mapped" can affect latency and have a substantial impact on your real-time system performance.